

Import packages to load data.

```
import pandas as pd # used for loading csv
import numpy as np # used to load and organize image data
```

Connect to drive.

```
from google.colab import drive
drive.mount("/content/drive")
```

Mounted at /content/drive

Train on Cuda if available:

```
import torch
train_on_gpu = torch.cuda.is_available()

if not train_on_gpu:
    print('CUDA is not available. Training on CPU ...')
else:
    print('CUDA is available! Training on GPU ...')
```

CUDA is not available. Training on CPU ...

Import data files.

```
images = np.load("/content/drive/MyDrive/Colab Notebooks/Advanced Analytics/images.npy", allow_pickle=True)
labels = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/Advanced Analytics/labels.csv")
```

The next steps are used to explore the data. We start by checking the shape of the images data and the head of the `labels` to make sure they have imported correctly.

```
print(images.shape)
print(labels.shape)
labels.head()
```

(4750, 128, 128, 3)  
(4750, 1)

Label



- 0 Small-flowered Cranesbill
- 1 Small-flowered Cranesbill
- 2 Small-flowered Cranesbill
- 3 Small-flowered Cranesbill
- 4 Small-flowered Cranesbill

Next steps: [Generate code with labels](#) [View recommended plots](#) [New interactive sheet](#)

Determine mean and standard deviation of images data to use in normalization later:

```
import torchvision.transforms as transforms

transform = transforms.Compose([
    transforms.ToTensor()])

images_tr = torch.stack([transform(image) for image in images])

sumel = 0.0
countel = 0
for img in images_tr:
    sumel += img.sum([1, 2])
    countel += torch.numel(img[0])
```

```
mean = sumel/countel
std = torch.sqrt(sumel/countel)
print("mean: " + str(mean))
print("standard deviation: " + str(std))

→ mean: tensor([0.2068, 0.2889, 0.3283])
    standard deviation: tensor([0.4548, 0.5375, 0.5730])
```

Get a list of plant labels being used.

```
labels['Label'].unique()

→ array(['Small-flowered Cranesbill', 'Fat Hen', 'Shepherds Purse',
       'Common wheat', 'Common Chickweed', 'Charlock', 'Cleavers',
       'Scentless Mayweed', 'Sugar beet', 'Maize', 'Black-grass',
       'Loose Silky-bent'], dtype=object)
```

Create list from categories.

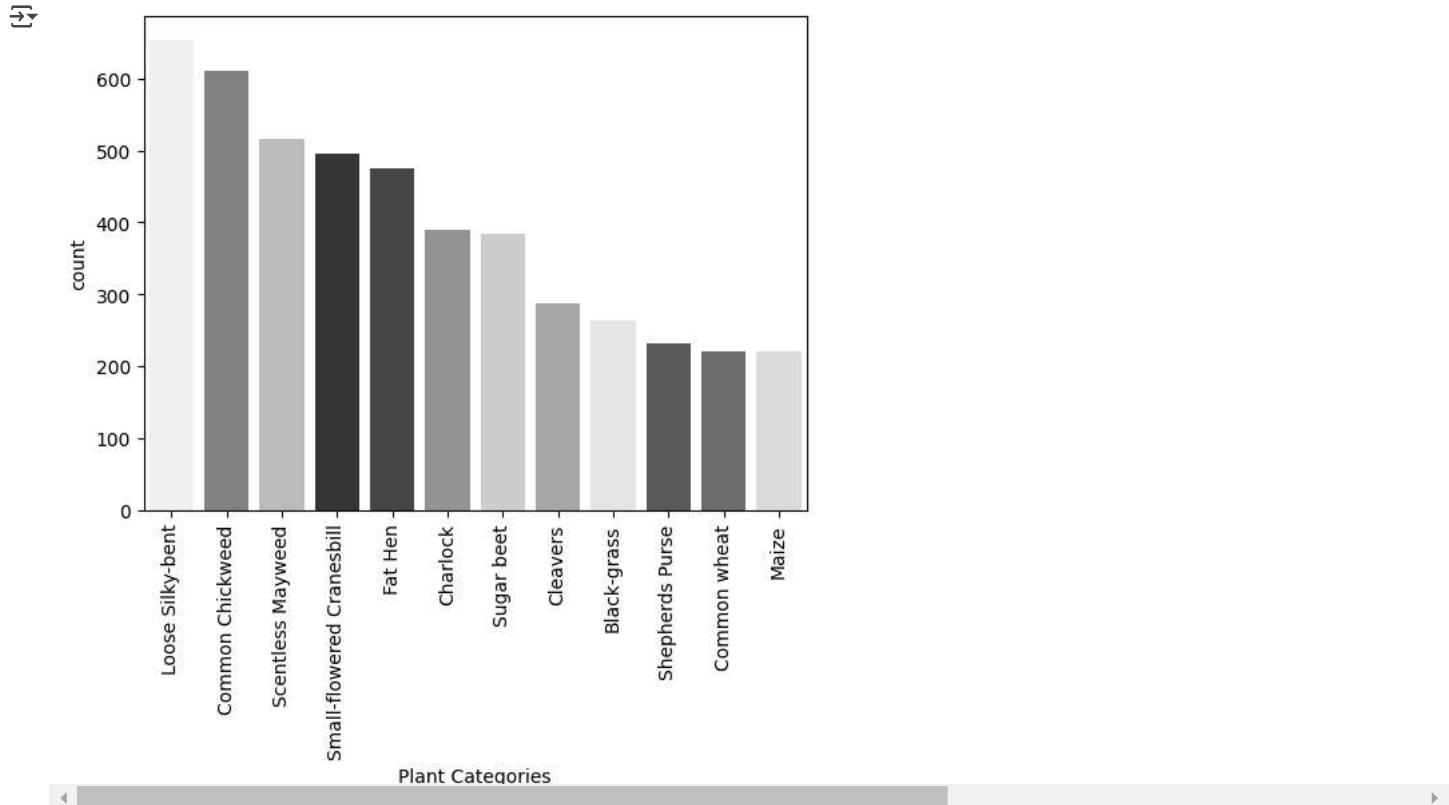
```
categories = ['Small-flowered Cranesbill', 'Fat Hen', 'Shepherds Purse',
              'Common wheat', 'Common Chickweed', 'Charlock', 'Cleavers',
              'Maize', 'Sugar beet', 'Scentless Mayweed', 'Black-grass',
              'Loose Silky-bent']
num_categories = len(categories) # will be used to create image grid
print(num_categories)
```

```
→ 12
```

Visualize the distribution of the categories.

```
import matplotlib.pyplot as plt # used to show plot and add information
import seaborn as sns # used to create plot

sns.countplot(x='Label', data=labels, order=labels['Label'].value_counts().index, legend=False, hue='Label', palette='Greens_r')
plt.xlabel('Plant Categories')
plt.xticks(rotation=90)
plt.show()
```



We can see that the Loose Silky-bent category contains more images than the others, and that common wheat and maize have the fewest images.

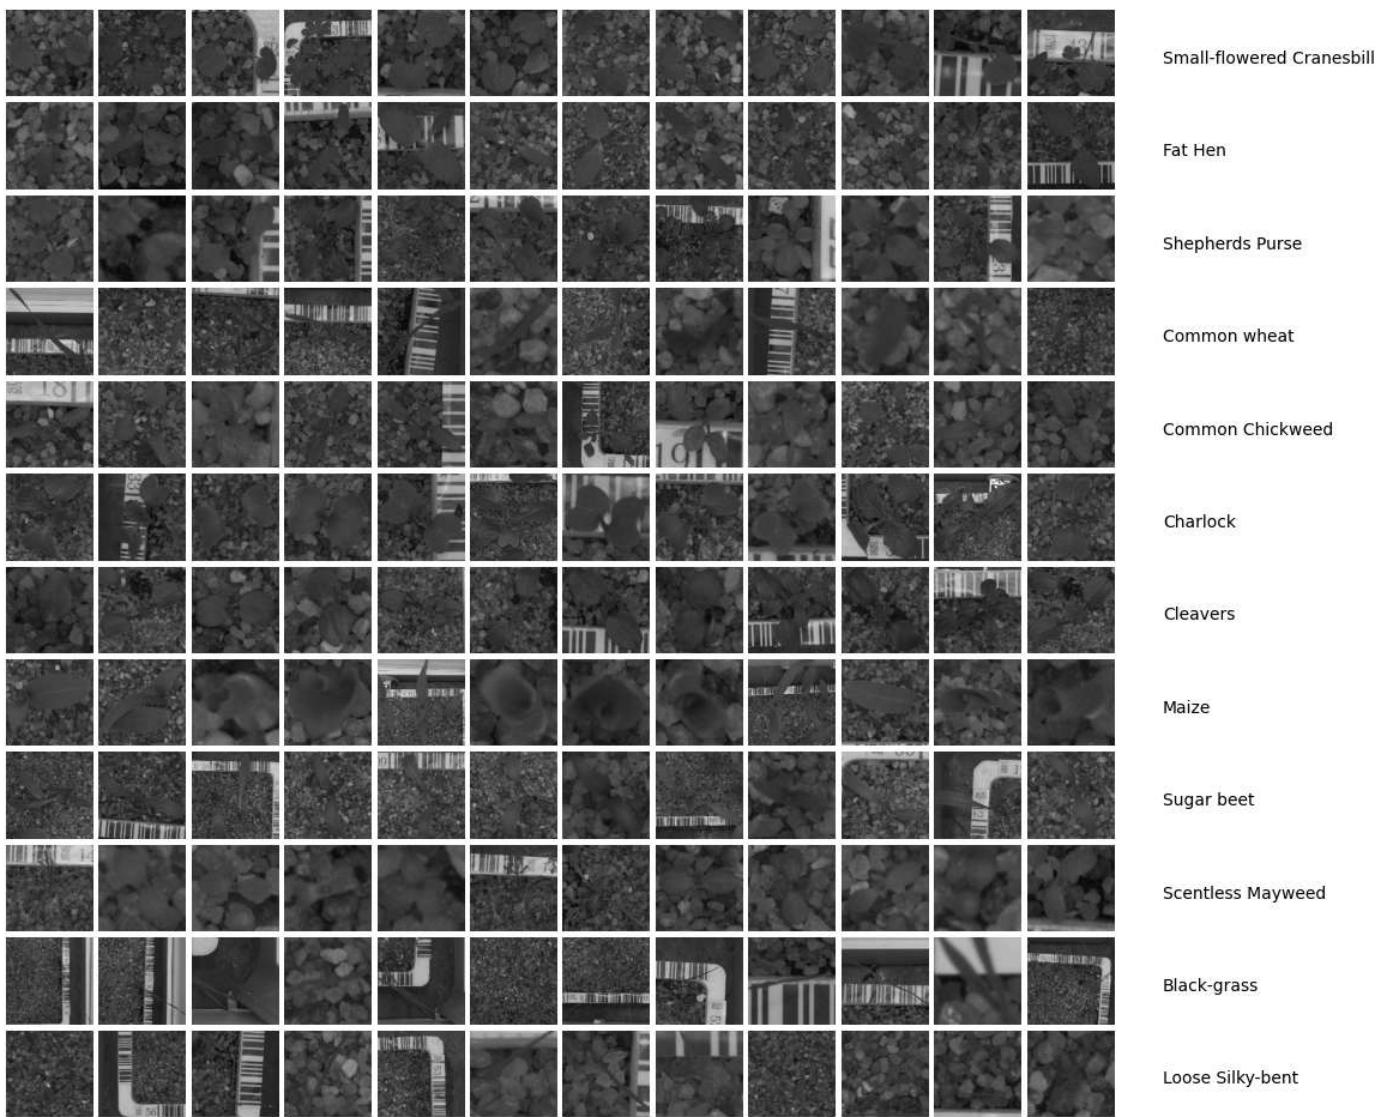
Visualize samples with associated categories:

```
from mpl_toolkits.axes_grid1 import ImageGrid # used to create image grid
import random
# Set seed for reproducibility
random_state = 12
random.seed(random_state)

# Set grid parameters
fig = plt.figure(1, figsize=(num_categories, num_categories))
grid = ImageGrid(fig, 111, nrows_ncols=(num_categories, num_categories), axes_pad=0.05)
i = 0

# Create loop for plotting images in each category
for category_id, category in enumerate(categories):
    condition = labels["Label"] == category
    plant_label = labels[condition].index.tolist()
    for k in range(min(12, len(plant_label))):
        ax = grid[i]
        ax.imshow(images[plant_label[k]])
        ax.axis('off')
        if i % num_categories == num_categories - 1:
            # used to print the names for each category
            ax.text(200, 70, category, verticalalignment='center')
        i += 1

plt.show()
```



One-hot encoding of categories:

```
import pandas as pd

print(labels.head(5))
labels = pd.get_dummies(labels, dtype=int)
#labels.reset_index(inplace=False)
print(labels.head(5))
print(labels.shape)

Label
0  Small-flowered Cranesbill
1  Small-flowered Cranesbill
2  Small-flowered Cranesbill
3  Small-flowered Cranesbill
4  Small-flowered Cranesbill
   Label_Black-grass  Label_Charlock  Label_Cleavers  Label_Common Chickweed  \
0            0            0            0            0
1            0            0            0            0
2            0            0            0            0
3            0            0            0            0
4            0            0            0            0

   Label_Common wheat  Label_Fat Hen  Label_Loose Silky-bent  Label_Maize  \
0            0            0            0            0
1            0            0            0            0
2            0            0            0            0
3            0            0            0            0
```

```

4          0          0          0          0
Label_Scentless Mayweed  Label_Shepherds Purse  \
0          0          0
1          0          0
2          0          0
3          0          0
4          0          0

Label_Small-flowered Cranesbill  Label_Sugar beet
0          1          0
1          1          0
2          1          0
3          1          0
4          1          0
(4750, 12)

```

Train-validation-test split:

```

from sklearn.model_selection import train_test_split

# Split into train and test (80%/20%)
X_train, X_test1, y_train, y_test1 = train_test_split(
    images, labels, test_size=0.20, stratify=labels, random_state = random_state)
# Split into test and validation (50/50 = 10%/10% of total)
X_val, X_test, y_val, y_test = train_test_split(
    X_test1, y_test1, test_size=0.50, stratify=y_test1, random_state = random_state)

print(X_train.shape, X_test.shape, y_train.shape, y_test.shape, X_val.shape, y_val.shape)

# Save data sets
from numpy import asarray
from numpy import savetxt
X_train_array = asarray(X_train)
#np.save('/content/drive/MyDrive/Colab Notebooks/Advanced Analytics/X_train.npy', X_train_array)
y_train_array = asarray(y_train)
#np.save('/content/drive/MyDrive/Colab Notebooks/Advanced Analytics/y_train.npy', y_train_array)
X_val_array = asarray(X_val)
#np.save('/content/drive/MyDrive/Colab Notebooks/Advanced Analytics/X_val.npy', X_val_array)
y_val_array = asarray(y_val)
#np.save('/content/drive/MyDrive/Colab Notebooks/Advanced Analytics/y_val.npy', y_val_array)
X_test_array = asarray(X_test)
#np.save('/content/drive/MyDrive/Colab Notebooks/Advanced Analytics/X_test.npy', X_test_array)
y_test_array = asarray(y_test)
#np.save('/content/drive/MyDrive/Colab Notebooks/Advanced Analytics/y_test.npy', y_test_array)

```

→ (3800, 128, 128, 3) (475, 128, 128, 3) (3800, 12) (475, 12) (475, 128, 128, 3) (475, 12)

```

# Define a custom dataset class to apply transforms to individual images
class CustomImageDataset(torch.utils.data.Dataset):
    def __init__(self, images, labels, transform):
        self.images = images
        self.labels = labels
        self.transform = transform

    def __getitem__(self, idx):
        image = self.images[idx]
        label = self.labels.iloc[idx].values
        if self.transform:
            image = self.transform(image)

        return image, label

    def __len__(self):
        return len(self.images)

```

Create customized dataset, transform training images to tensor, normalize, and preview one batch:

```

from torchvision import datasets
import torchvision.transforms as transforms
import torch
from torch.utils.data import DataLoader, TensorDataset
import multiprocessing
import matplotlib.pyplot as plt
%matplotlib inline

```

```

from torchvision.datasets import ImageFolder
from PIL import Image

n_workers = multiprocessing.cpu_count()
print(n_workers)

mean = [0.2068, 0.2889, 0.3283]
std = [0.4548, 0.5375, 0.5730]

train_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean, std)
])

# Create instance of the custom dataset
train_data = CustomImageDataset(images=X_train, labels=y_train, transform=train_transforms)

# Create DataLoader
train_loader = DataLoader(train_data, batch_size=20, num_workers=n_workers, shuffle=True)

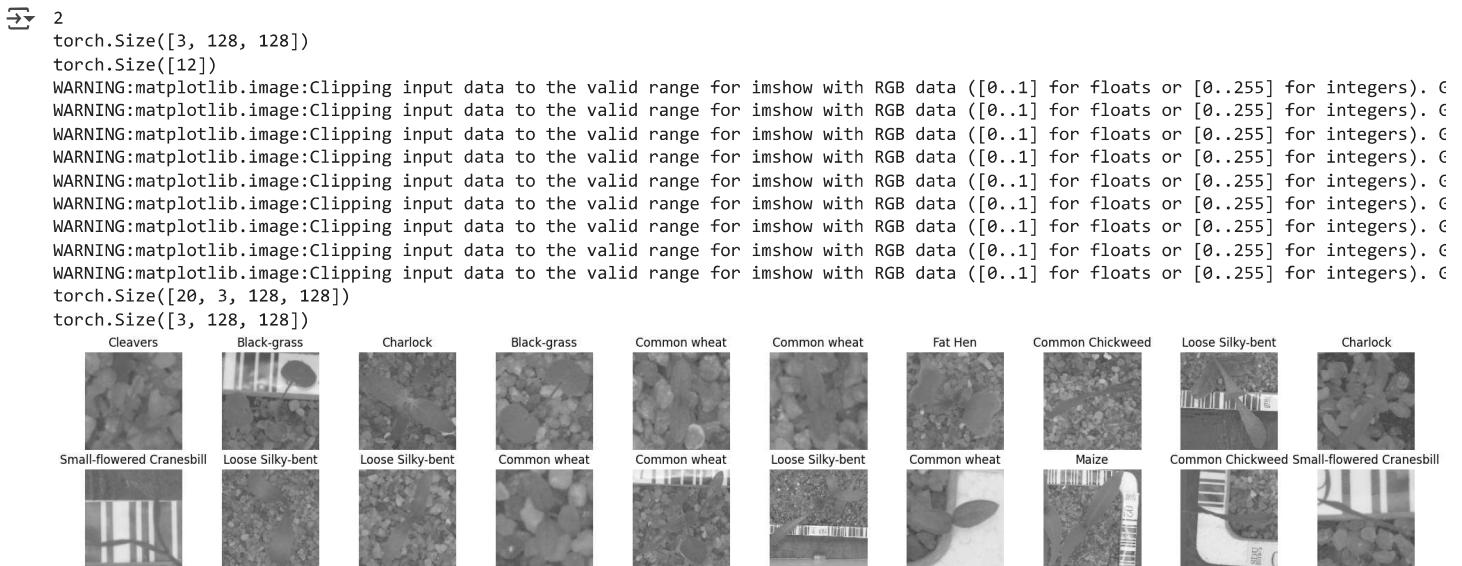
# Obtain next batch
images, labels = next(iter(train_loader))
image = images[0]
label = labels[0]
print(image.shape)
print(label.shape)

# Function to display image
def imshow(img, sub):
    img = img / 2 + 0.5
    img = img.numpy()
    img = img.transpose((1, 2, 0))
    sub.imshow(img)
    sub.axis('off')

# Plot images
fig, subs = plt.subplots(2, 10, figsize=(25, 4))
for i, sub in enumerate(subs.flatten()):
    imshow(images[i], sub)
    # Get the index of the maximum value instead of the column name
    category_index = labels[i].argmax()
    sub.set_title(categories[category_index])

print(images.shape)
print(image.shape)

```



Create function for transforms with augmentations and create custom dataset and dataloaders.

```

from torchvision import datasets
import torchvision.transforms as transforms
import torch

```

```
from torch.utils.data import DataLoader, TensorDataset
import multiprocessing
import matplotlib.pyplot as plt
%matplotlib inline
from torchvision.datasets import ImageFolder
from PIL import Image
import numpy as np

n_workers = multiprocessing.cpu_count()
print(n_workers)

# Create function for transformations for each dataset
# Function will allow for optimization of parameters
def get_transforms():
    mean = [0.2068, 0.2889, 0.3283]
    std = [0.4548, 0.5375, 0.5730]

    return{
        "train": transforms.Compose([
            transforms.RandomHorizontalFlip(),
            transforms.RandomVerticalFlip(),
            transforms.RandomRotation(45),
            transforms.RandomAutocontrast(),
            transforms.ToTensor(),
            transforms.Normalize(mean, std)
        ]),
        'validate': transforms.Compose([
            transforms.ToTensor(),
            transforms.Normalize(mean, std)
        ]),
        'test': transforms.Compose([
            transforms.ToTensor(),
            transforms.Normalize(mean, std)
        ])
    }

# Define a custom dataset class to apply transforms to individual images
class CustomImageDataset(torch.utils.data.Dataset):
    def __init__(self, images, labels, transform=None):
        self.images = images
        self.labels = labels
        self.transform = transform

    def __len__(self):
        return len(self.images)

    def __getitem__(self, idx):
        image = self.images[idx]
        label = self.labels.iloc[idx].values # Assuming labels is a pandas DataFrame

        # Convert the NumPy array to a PIL Image for augmentation
        image = Image.fromarray(image)

        if self.transform:
            image = self.transform(image)

        return image, label

#Create function for data loaders
def get_data_loaders(transforms, num_workers, random_seed=12):
    # Reseed random number generators for comparing experiments
    torch.manual_seed(random_seed)
    torch.cuda.manual_seed(random_seed)
    np.random.seed(random_seed)

    transforms = get_transforms()
    # Create each instance of the custom dataset
    train_data = CustomImageDataset(images=X_train, labels=y_train, transform=transforms['train'])
    val_data = CustomImageDataset(images=X_val, labels=y_val, transform=transforms['validate'])
    test_data = CustomImageDataset(images=X_test, labels=y_test, transform=transforms['test'])

    # Create data loaders
    train_loader = DataLoader(train_data, batch_size=100, num_workers=n_workers, shuffle=True)
    val_loader = DataLoader(val_data, batch_size=25, num_workers=n_workers, shuffle=True)
    test_loader = DataLoader(test_data, batch_size=25, num_workers=n_workers, shuffle=True)

    return {'train': train_loader, 'validate': val_loader, 'test': test_loader}
```

→ 2

Define model structure:

```

import torch.nn as nn
from torchsummary import summary

# define CNN model
class Net(nn.Module):
    def __init__(self, num_classes=12):
        super(Net, self).__init__()

        self.model = nn.Sequential(
            # convolutional layer 1 (tensor 3x128x128 -> 16x64x64)
            nn.Conv2d(in_channels=3, out_channels=16, kernel_size=3, padding=1),
            nn.BatchNorm2d(16),
            nn.ELU(),
            nn.MaxPool2d(kernel_size=2, stride=2),
            # convolutional layer 2 - takes output of previous layer (16x64x64 -> 32x32x32)
            nn.Conv2d(16, 32, 3, padding=1),
            nn.BatchNorm2d(32),
            nn.ELU(),
            nn.MaxPool2d(2, 2),
            # convolutional layer 3 (32x32x32 -> 164x6x16)
            nn.Conv2d(32, 64, 3, padding=1),
            nn.BatchNorm2d(64),
            nn.ELU(),
            nn.MaxPool2d(2, 2),
            # linear layer (64x16x16 -> 500)
            nn.Flatten(),
            nn.Linear(64 * 16 * 16, 500),
            nn.Dropout(0.5),
            nn.BatchNorm1d(500),
            nn.ELU(),
            # linear layer (500 -> 12)
            nn.Linear(500, num_classes)
        )

    def forward(self, x):
        return self.model(x)

# create complete CNN
model = Net()
print(model)
summary(model, input_size=(3, 128, 128))

→ Net(
  (model): Sequential(
    (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ELU(alpha=1.0)
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (4): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (5): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (6): ELU(alpha=1.0)
    (7): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (8): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (10): ELU(alpha=1.0)
    (11): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (12): Flatten(start_dim=1, end_dim=-1)
    (13): Linear(in_features=16384, out_features=500, bias=True)
    (14): Dropout(p=0.5, inplace=False)
    (15): BatchNorm1d(500, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (16): ELU(alpha=1.0)
    (17): Linear(in_features=500, out_features=12, bias=True)
  )
)
-----
      Layer (type)          Output Shape       Param #
=====
      Conv2d-1           [-1, 16, 128, 128]      448
      BatchNorm2d-2      [-1, 16, 128, 128]      32
      ELU-3              [-1, 16, 128, 128]      0
      MaxPool2d-4        [-1, 16, 64, 64]        0
      Conv2d-5           [-1, 32, 64, 64]      4,640
      BatchNorm2d-6      [-1, 32, 64, 64]        64
      ELU-7              [-1, 32, 64, 64]        0

```

```
MaxPool2d-8      [-1, 32, 32, 32]          0
Conv2d-9        [-1, 64, 32, 32]         18,496
BatchNorm2d-10   [-1, 64, 32, 32]         128
    ELU-11       [-1, 64, 32, 32]          0
MaxPool2d-12     [-1, 64, 16, 16]          0
    Flatten-13    [-1, 16384]             0
    Linear-14     [-1, 500]              8,192,500
    Dropout-15    [-1, 500]             1,000
BatchNorm1d-16   [-1, 500]             1,000
    ELU-17       [-1, 500]             1,000
    Linear-18     [-1, 12]              6,012
=====
Total params: 8,223,320
Trainable params: 8,223,320
Non-trainable params: 0
-----
Input size (MB): 0.19
Forward/backward pass size (MB): 11.52
Params size (MB): 31.37
Estimated Total Size (MB): 43.07
-----
```

Create functions for training loop and validation/test loop:

```
!pip install livelossplot
```

→ Show hidden output

```
!pip install torchmetrics
```

```
→ Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (883 kB)
    24.6/24.6 MB 26.0 MB/s eta 0:00:00
    883.7/883.7 kB 32.1 MB/s eta 0:00:00
    Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl (664.8 MB)
```

```
Successfully uninstalled nvidia-cuda-cu12-11.6.3.83
Attempting uninstall: nvidia-cusolver-cu12
Found existing installation: nvidia-cusolver-cu12 11.6.3.83
Uninstalling nvidia-cusolver-cu12-11.6.3.83:
Successfully uninstalled nvidia-cusolver-cu12-11.6.3.83
Successfully installed lightning-utilities-0.12.0 nvidia-cublas-cu12-12.4.5.8 nvidia-cuda-cupti-cu12-12.4.127 nvidia-cuda-nvrtc-cu12-12.4.127

import torch.optim as optim
import liveLossPlot
from liveLossPlot import PlotLosses
from torchmetrics import Precision, Recall

# train on cuda if available
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')

# Define the loss function
criterion = nn.CrossEntropyLoss()

def train_one_epoch(data_loaders, model, optimizer, criterion):
    liveLoss = PlotLosses()
    for epoch in range(1):
        model.train()
        running_loss = 0.0
        running_corrects = 0
        logs = {}
        # Iterate over training batches
        for images, labels in data_loaders:
            optimizer.zero_grad()
            outputs = model(images)
            labels = torch.argmax(labels, dim=1)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()
            running_loss += loss.item()

        train_loss = running_loss / len(data_loaders)
        logs['train log loss'] = train_loss
        print(f"Training loss: {train_loss:.3f}")
        return train_loss
    liveLoss.update(logs)
    liveLoss.send()

def val_one_epoch(data_loaders, model, criterion):
    liveLoss = PlotLosses()
    with torch.no_grad():
        for epoch in range(1):
            model.eval()
            running_loss = 0.0
            running_corrects = 0
            logs = {}
            # Initialize lists for storing predictions and labels
            preds = []
            actuals = []
            # Iterate over training batches
            for images, labels in data_loaders:
                outputs = model(images)
                labels = torch.argmax(labels, dim=1)
                loss = criterion(outputs, labels)
                running_loss += loss.item()
                _, predicted = torch.max(outputs, 1)
                running_corrects += torch.sum(predicted == labels.data)

            val_loss = running_loss / len(data_loaders)
            val_accuracy = float(running_corrects) / len(data_loaders.dataset)
            logs['val log loss'] = val_loss
            logs['val accuracy'] = val_accuracy
            print(f"Validation loss - validate: {val_loss:.3f}")
            print(f"Validation Accuracy: {val_accuracy:.3f}")
            return val_loss, val_accuracy
    liveLoss.update(logs)
    liveLoss.send()

def test_one_epoch(data_loaders, model, criterion):
    with torch.no_grad():
        for epoch in range(1):
            # Define metrics
            metric_precision = Precision(task = 'multiclass', num_classes = 12, average = None)
```

```

metric_recall = Recall(task = 'multiclass', num_classes = 12, average = None)
model.eval()
running_loss = 0
running_corrects = 0
# Initialize lists for storing predictions and labels
preds = []
actuals = []
# Iterate over training batches
for images, labels in data_loaders:
    outputs = model(images)
    labels = torch.argmax(labels, dim=1)
    loss = criterion(outputs, labels)
    running_loss += loss.item()

    # Get and save predicted labels
    _, predicted = torch.max(outputs, 1)
    running_corrects += torch.sum(predicted == labels.data)
    preds.extend(outputs.cpu().numpy())
    actuals.extend(labels.cpu().numpy())

# Convert predictions and actuals to numpy array
preds = torch.tensor(np.array(preds))
actuals = torch.tensor(np.array(actuals))

# Calculate metrics
metric_precision(preds, actuals)
metric_recall(preds, actuals)

val_loss = running_loss / len(data_loaders.dataset)
accuracy = float(running_corrects) / len(data_loaders.dataset)
print(f"Validation loss - test: {val_loss:.3f}")
print(f"Test Accuracy: {accuracy:.3f}")
precision = metric_precision.compute()
recall = metric_recall.compute()
print(f"Precision: {precision}")
print(f"Recall: {recall}")
return val_loss, preds, actuals

```

Create early stopping class:

```

class EarlyStopper:
    def __init__(self, patience=1, min_delta=0):
        self.patience = patience
        self.min_delta = min_delta
        self.counter = 0
        self.min_val_loss = float('inf')

    def early_stop(self, val_loss):
        if val_loss < self.min_val_loss:
            self.min_val_loss = val_loss
            self.counter = 0
        elif val_loss > (self.min_val_loss + self.min_delta):
            self.counter += 1
            if self.counter >= self.patience:
                return True
        return False

```

Define train one model:

```
!pip install mlflow
```

Show hidden output

```

import mlflow
from torch.optim.lr_scheduler import ExponentialLR
import datetime # Used for experiment name
import os

# train on cuda if available
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
n_workers = multiprocessing.cpu_count()
save_path = "/content/drive/MyDrive/Colab Notebooks/Advanced Analytics"

```

```

def train_one_model(learning_rate, gamma, n_epochs):
    liveloss = PlotLosses()
    transforms = get_transforms()
    data_loaders = get_data_loaders(transforms, num_workers=n_workers)
    model = Net()
    val_loss_min = None

    if train_on_gpu:
        model.cuda()

    optimizer = optim.Adam(model.parameters(), lr = learning_rate)
    scheduler = ExponentialLR(optimizer, gamma=gamma)
    criterion = nn.CrossEntropyLoss()

    with mlflow.start_run():
        for epoch in range(n_epochs):
            logs= {}
            print("Epoch: ", epoch)
            train_loss = train_one_epoch(data_loaders['train'], model, optimizer, criterion)

            val_loss, val_accuracy = val_one_epoch(data_loaders['validate'], model, criterion)

            # Include early stop if necessary
            early_stopper = EarlyStopper(patience=3)
            if early_stopper.early_stop(val_loss):
                print("We are at epoch:", epoch)
                break

            # Save model if validation loss decreases by more than 1%
            if val_loss_min is None or
                (val_loss_min - val_loss) / val_loss_min > 0.01
            ):

                # Save weights
                filename = os.path.join(save_path, 'best_val_loss.pt')
                torch.save(model.state_dict(), filename)

                val_loss_min = val_loss

            logs['validation loss'] = val_loss
            logs['training loss'] = train_loss
            logs["validation accuracy"] = val_accuracy
            liveloss.update(logs)
            liveloss.send()

            # Update learning rate
            scheduler.step(val_loss)

            # Test model on validation set (no optimization)
            val_loss, preds, actuals = test_one_epoch(data_loaders['validate'], model, criterion)

            # Restore best validation loss parameter
            model.load_state_dict(torch.load(filename))

            # Log hyperparameters
            mlflow.log_metric("val_loss", val_loss)
            mlflow.log_param("learning_rate", learning_rate)
            mlflow.log_param("n_epochs", n_epochs)
            mlflow.log_param("gamma", gamma)

            # Log metrics
            mlflow.log_metric("train_loss", train_loss)
            mlflow.log_metric("val_loss", val_loss)
            val_accuracy = (np.argmax(preds, axis=1) == np.array(actuals)).sum() / len(actuals)
            mlflow.log_metric("val_accuracy", val_accuracy)
            mlflow.log_artifact(filename)

        return model
    mlflow.end_run()

```

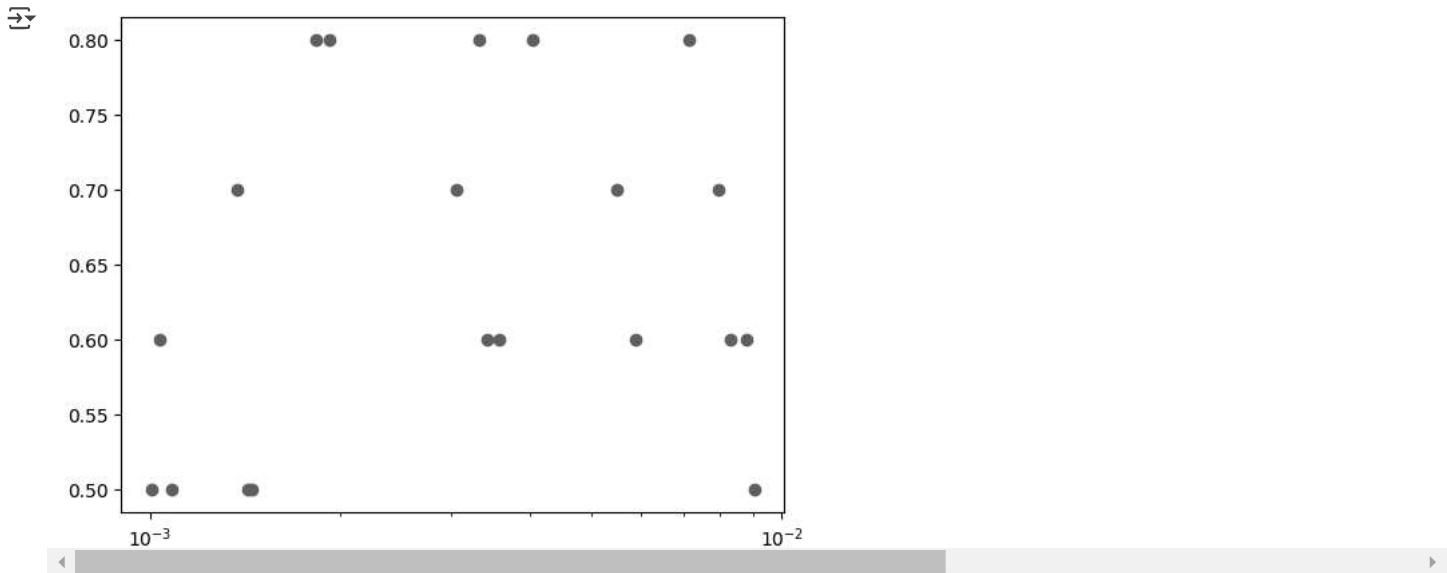
Use random search to explore hyperparameters:

```
# Create random grid of learning rate and gamma
min_lr = 0.001
```

```
max_lr = 0.01
n_grid = 20

# Sample learning rate log samples
lrs = 10**np.random.uniform(np.log10(min_lr), np.log10(max_lr), n_grid))
# Uniformly sample gamma rate
gamma = np.random.randint(5, 9, n_grid) / 10

# Plot grid
_ = plt.scatter(lrs, gamma)
_ = plt.xscale("log")
```



Run experiment to determine best learning-rate and gamma values:

```
for lr, g in zip(lrs, gamma):
    train_one_model(lr, g, n_epochs=5)

Epoch: 0
Training loss: 1.779
Validation loss - validate: 1.661
Validation Accuracy: 0.453
Validation loss - test: 0.066
Test Accuracy: 0.453
Epoch: 1
<ipython-input-25-4c4b3fabe1a4>:63: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default value),
    model.load_state_dict(torch.load(filename))
Training loss: 1.284
Validation loss - validate: 1.212
Validation Accuracy: 0.615
Validation loss - test: 0.048
Test Accuracy: 0.615
Epoch: 2
Training loss: 1.082
Validation loss - validate: 1.042
Validation Accuracy: 0.642
Validation loss - test: 0.042
Test Accuracy: 0.642
Epoch: 3
Training loss: 0.998
Validation loss - validate: 1.001
Validation Accuracy: 0.651
Validation loss - test: 0.040
Test Accuracy: 0.651
Epoch: 4
Training loss: 0.917
Validation loss - validate: 0.947
Validation Accuracy: 0.655
Validation loss - test: 0.038
Test Accuracy: 0.655
Epoch: 0
Training loss: 1.749
Validation loss - validate: 1.534
Validation Accuracy: 0.495
Validation loss - test: 0.061
Test Accuracy: 0.495
Epoch: 1
```

```

Training loss: 1.241
Validation loss - validate: 1.093
Validation Accuracy: 0.625
Validation loss - test: 0.044
Test Accuracy: 0.625
Epoch: 2
Training loss: 1.043
Validation loss - validate: 1.128
Validation Accuracy: 0.640
Validation loss - test: 0.045
Test Accuracy: 0.640
Epoch: 3
Training loss: 1.075
Validation loss - validate: 1.048
Validation Accuracy: 0.632
Validation loss - test: 0.042
Test Accuracy: 0.632
Epoch: 4

```

Retrieve mlflow data on experiment:

```

import mlflow

runs = mlflow.search_runs()
sorted_runs = runs[
    ["run_id",
     "params.learning_rate",
     "params.n_epochs",
     "params.gamma",
     "metrics.train_loss",
     "metrics.val_loss",
     "metrics.val_accuracy"
    ]
].sort_values(by="metrics.val_loss", ascending=True)
sorted_runs

```

	run_id	params.learning_rate	params.n_epochs	params.gamma	metrics.train_loss	metrics.val_loss	metric
14	a99b773d328a47d892da69bdea79e363	0.008293674908605475	5	0.6	0.929735	0.032949	
1	066614e2ad7b4c12be4fbead33459323	0.005863192512339714	5	0.6	0.884745	0.033293	
2	a86aeb3916f845e8bdcb24ab3d28ab75	0.003057577076255498	5	0.7	0.898287	0.033676	
7	c21ee1cd95924147acc3b554725c532	0.008794783157695697	5	0.6	0.912329	0.033954	
6	943325aba9fe4d498495ccecd8ab436e	0.007124190781828619	5	0.8	0.893025	0.034535	
16	59ad864641d64db9a0256d2ff0f9a13e	0.0034177429283790495	5	0.6	0.892751	0.035029	
18	b895821e6f4a438389374b937c4f2020	0.005496037616682325	5	0.7	0.950851	0.035216	
4	28c21eab661c4d7ca66cedb8dad2cb7d	0.0033206723595393755	5	0.8	0.895180	0.035251	
17	e2569cf3b43a4428b0b420a613251613	0.0018336439724345017	5	0.8	0.908110	0.035508	
8	e24e6c2bb94d4fab08583b9cc674a43	0.004037227142333932	5	0.8	0.884938	0.036171	
12	e68df2efea4e48fb8d69331c6e560805	0.001079994211106168	5	0.5	0.931614	0.036356	
15	20538e0457f54f7b66422b962c2c7ce	0.001034129584101865	5	0.6	0.932521	0.036451	
0	056fcf804e1a4c2e93ac9818bd71746c	0.0014478272725587603	5	0.5	0.916393	0.036469	
9	7097b88a839c4e69b39d9ba4bf18e607	0.0019223318118548767	5	0.8	0.906705	0.036512	
5	6f3c97aaaa9342f695bea08d93ea680d	0.0010052156317193804	5	0.5	0.935427	0.036734	
3	27e02fece26b408ebef8bbfd584786f	0.003564820226076412	5	0.6	0.887444	0.036880	
10	f4527b932658420db3ad905570d5aa5d	0.0013715426633363898	5	0.7	0.916864	0.037160	
19	4052cdff09d64ea0ae9862ae6c5375d1	0.0014261422374240921	5	0.5	0.916973	0.037865	
11	e015346caffa414fb44137d40ce6704a	0.00905626946526	5	0.5	1.077178	0.039753	
13	69d740a70be24e09bd866b5f7e79cc09	0.00795636785524916	5	0.7	1.071292	0.041005	
20	fbfb147191d944c978a9a56a81005f7ed	0.0014261422374240921	10	0.5	1.081689	0.041686	
21	1134b37b3f094aae9baeb0adbcc98128	None	None	None	NaN	NaN	

Next steps: [Generate code with sorted\\_runs](#) [View recommended plots](#) [New interactive sheet](#)

Save model with best parameters in case I need it later:

```
from mlflow.tracking import MlflowClient

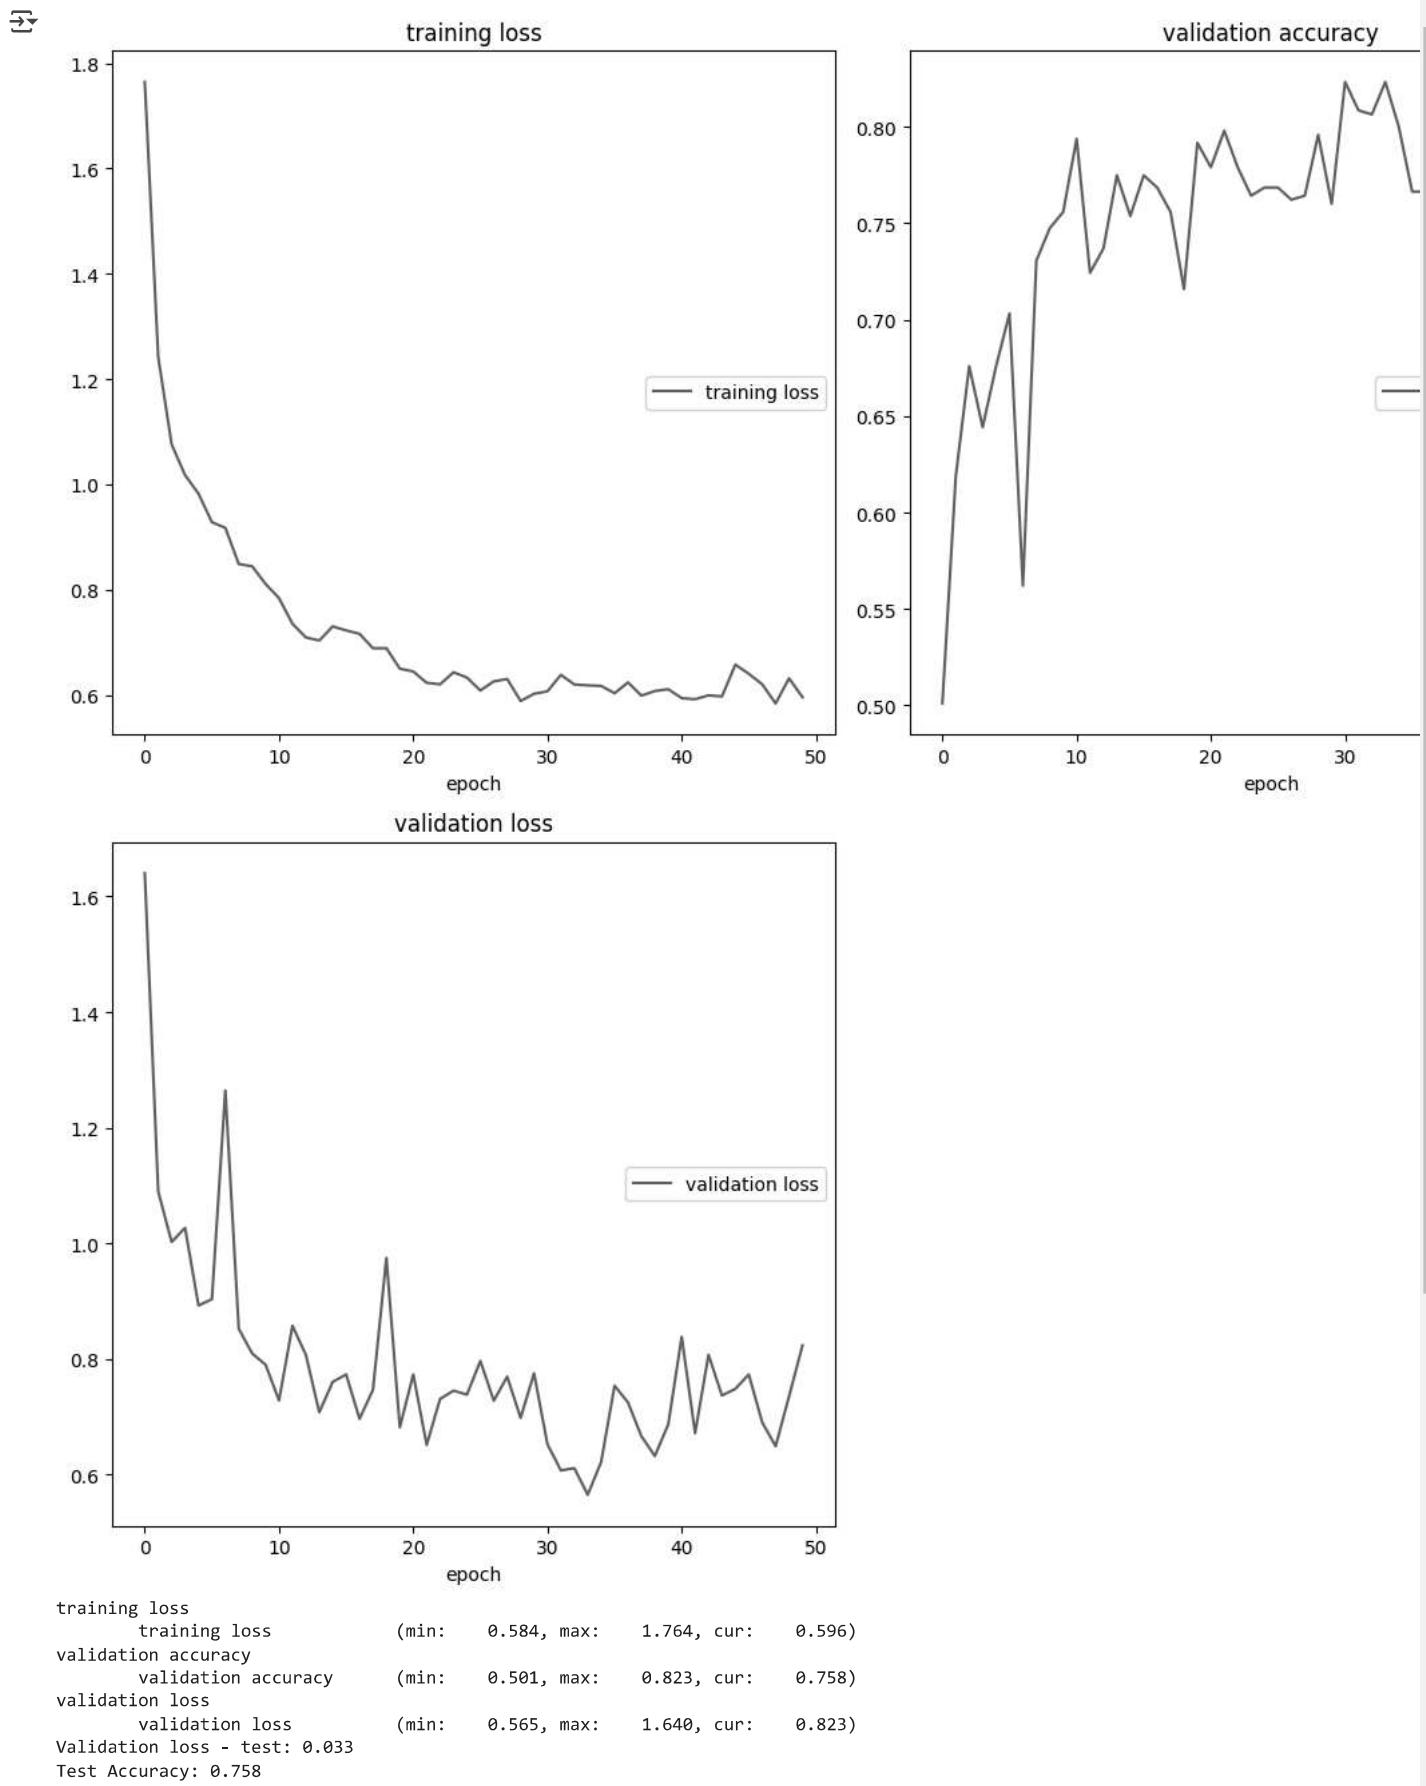
lowest_loss_id = sorted_runs.iloc[0]['run_id']

# Fetch model for that run
client = MlflowClient()
local_path = client.download_artifacts(lowest_loss_id, "best_val_loss.pt", '.')


```

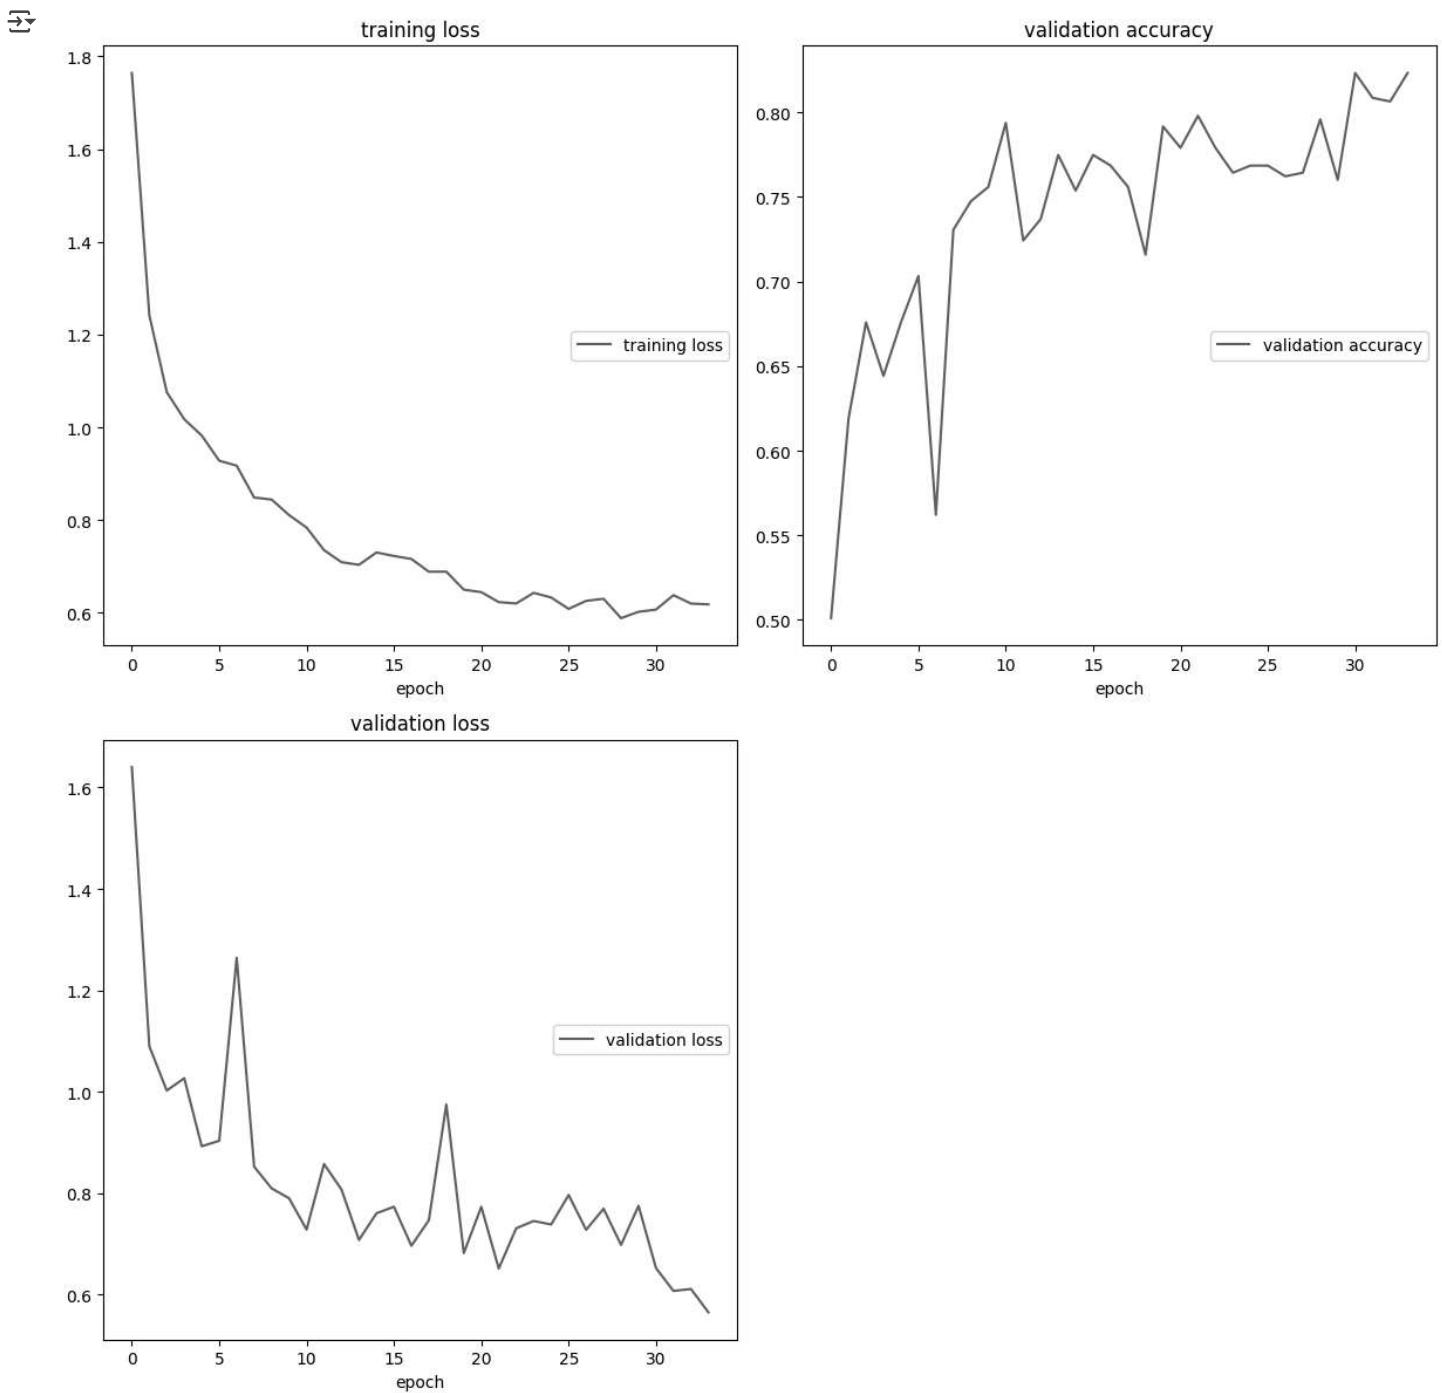
Train model with selected parameters to determine optimal number of epochs to use:

```
model_train = train_one_model(learning_rate=0.008, gamma=0.6, n_epochs=50)
```



Based on this experiment, using 34 epochs will be an ideal number. I will train a final model using these parameters.

```
final_model = train_one_model(learning_rate=0.008, gamma=0.6, n_epochs=34)
```



```

training loss
    training loss      (min: 0.589, max: 1.764, cur: 0.618)
validation accuracy
    validation accuracy (min: 0.501, max: 0.823, cur: 0.823)
validation loss
    validation loss     (min: 0.565, max: 1.640, cur: 0.565)
Validation loss - test: 0.023
Test Accuracy: 0.823
Precision: tensor([0.4643, 0.8750, 0.9545, 0.9091, 0.7692, 0.8235, 0.7719, 0.9444, 0.7288,
0.7083, 0.9574, 0.9189])
Recall: tensor([0.5000, 0.8974, 0.7241, 0.9836, 0.9091, 0.8750, 0.6769, 0.7727, 0.8431,
0.7391, 0.9000, 0.8718])

```

Save final model:

```

model_scripted = torch.jit.script(final_model) # Export to TorchScript
model_scripted.save('model_scripted.pt') # Save

```